

mini-project-2

April 12, 2024

1 Project Overview

This notebook presents the development of a modified ResNet architecture aimed at achieving high classification accuracy on the CIFAR-10 dataset, with the constraint of maintaining fewer than 5 million parameters. The project explores various architectural, optimization, and data processing strategies to meet these objectives.

```
[1]: import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
from torch.utils.data import DataLoader
import numpy as np
import random
import matplotlib.pyplot as plt

SEED = 42
random.seed(SEED)
np.random.seed(SEED)
torch.manual_seed(SEED)
torch.cuda.manual_seed(SEED)
torch.backends.cudnn.deterministic = True

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
```

```
[2]: def evaluate(model, data_loader):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in data_loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
```

```

accuracy = correct / total
return accuracy

```

```

[3]: transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ColorJitter(brightness=0.1, contrast=0.1, saturation=0.1, hue=0.
↪1),
    transforms.RandomRotation(10),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
↪download=True, transform=transform_train)
train_loader = DataLoader(trainset, batch_size=128, shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
↪download=True, transform=transform_test)
test_loader = DataLoader(testset, batch_size=100, shuffle=False)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse',
↪'ship', 'truck']

```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to
./data/cifar-10-python.tar.gz

100%| | 170498071/170498071 [00:06<00:00, 28042028.38it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

```

[4]: class BasicBlock(nn.Module):
    expansion = 1

    def __init__(self, in_channels, out_channels, stride=1):
        super(BasicBlock, self).__init__()
        self.conv1 = nn.Conv2d(in_channels, out_channels, kernel_size=3,
↪stride=stride, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(out_channels)
        self.conv2 = nn.Conv2d(out_channels, out_channels, kernel_size=3,
↪stride=1, padding=1, bias=False)

```

```

        self.bn2 = nn.BatchNorm2d(out_channels)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_channels != self.expansion * out_channels:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_channels, self.expansion * out_channels,
↪kernel_size=1, stride=stride, bias=False),
                nn.BatchNorm2d(self.expansion * out_channels)
            )

        def forward(self, x):
            out = torch.relu(self.bn1(self.conv1(x)))
            out = self.bn2(self.conv2(out))
            out += self.shortcut(x)
            out = torch.relu(out)
            return out

class ResNet(nn.Module):
    def __init__(self, block, num_blocks, num_classes=10):
        super(ResNet, self).__init__()
        self.in_planes = 66

        self.conv1 = nn.Conv2d(3, 66, kernel_size=3, stride=1, padding=1,
↪bias=False)
        self.bn1 = nn.BatchNorm2d(66)

        self.layer1 = self._make_layer(block, 66, num_blocks[0], stride=1)

        self.layer2 = self._make_layer(block, 132, num_blocks[1], stride=2)
        self.layer3 = self._make_layer(block, 264, num_blocks[2], stride=2)

        self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
        self.linear = nn.Linear(264 * block.expansion, num_classes)

    def _make_layer(self, block, planes, num_blocks, stride):
        strides = [stride] + [1]*(num_blocks-1)
        layers = []
        for stride in strides:
            layers.append(block(self.in_planes, planes, stride))
            self.in_planes = planes * block.expansion
        return nn.Sequential(*layers)

    def forward(self, x):
        out = torch.relu(self.bn1(self.conv1(x)))
        out = self.layer1(out)
        out = self.layer2(out)

```

```

        out = self.layer3(out)
        out = self.avgpool(out)
        out = out.view(out.size(0), -1)
        out = self.linear(out)
        return out

def ResNetCustom():
    return ResNet(BasicBlock, [3, 4, 3], num_classes=10)

model = ResNetCustom().to(device)

```

```

[5]: import torch
import torch.nn as nn
import torch.nn.functional as F

class SmoothCrossEntropyLoss(nn.Module):
    def __init__(self, smoothing=0.1):
        super(SmoothCrossEntropyLoss, self).__init__()
        self.smoothing = smoothing

    def forward(self, input, target):
        log_prob = F.log_softmax(input, dim=-1)
        weight = input.new_ones(input.size()) * self.smoothing / (input.
↪size(-1) - 1.)
        weight.scatter_(-1, target.unsqueeze(-1), (1. - self.smoothing))
        loss = (-weight * log_prob).sum(dim=-1).mean()
        return loss

```

```

[6]: from torch.optim.lr_scheduler import LambdaLR
import math
model = ResNetCustom().to(device)
criterion = SmoothCrossEntropyLoss(smoothing=0.1).to(device)
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-4)

def warmup_with_cosine_lr(epoch):
    if epoch < 5:
        return epoch / 5
    else:
        return 0.5 * (1. + math.cos(math.pi * (epoch - 5) / (epochs - 5)))

scheduler = LambdaLR(optimizer, lr_lambda=warmup_with_cosine_lr)

```

```

[7]: model_parameters = sum(p.numel() for p in model.parameters() if p.requires_grad)
print(f"Number of Trainable Parameters: {model_parameters}")

```

Number of Trainable Parameters: 4916284

```
[8]: train_losses, val_losses, train_accuracies, val_accuracies = [], [], [], []
```

```
def train(epoch):
    model.train()
    train_loss, correct, total = 0, 0, 0
    for inputs, targets in train_loader:
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

        train_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()

    train_losses.append(train_loss / len(train_loader))
    train_accuracies.append(100. * correct / total)
    print(f'Epoch: {epoch}, Train Loss: {train_losses[-1]:.4f}, Train Acc:␣
↪{train_accuracies[-1]:.2f}%')

def validate(epoch):
    model.eval()
    val_loss, correct, total = 0, 0, 0
    with torch.no_grad():
        for inputs, targets in test_loader:
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = model(inputs)
            loss = criterion(outputs, targets)

            val_loss += loss.item()
            _, predicted = outputs.max(1)
            total += targets.size(0)
            correct += predicted.eq(targets).sum().item()

    val_losses.append(val_loss / len(test_loader))
    val_accuracies.append(100. * correct / total)
    print(f'Epoch: {epoch}, Val Loss: {val_losses[-1]:.4f}, Val Acc:␣
↪{val_accuracies[-1]:.2f}%')
```

```
[9]: epochs = 60
train_losses, val_losses, train_accuracies, val_accuracies = [], [], [], []
lr_changes = []

for epoch in range(epochs):
```

```

train(epoch)
validate(epoch)
scheduler.step()
lr_changes.append(optimizer.param_groups[0]['lr'])

model_save_path = './model_weights.pth'
torch.save(model.state_dict(), model_save_path)
print(f'Model saved to {model_save_path}')

```

/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork() was called. os.fork() is incompatible with multithreaded code, and JAX is multithreaded, so this will likely lead to a deadlock.

```

self.pid = os.fork()

```

```

Epoch: 0, Train Loss: 2.4132, Train Acc: 10.76%
Epoch: 0, Val Loss: 2.3967, Val Acc: 11.08%
Epoch: 1, Train Loss: 1.6181, Train Acc: 49.15%
Epoch: 1, Val Loss: 1.5646, Val Acc: 56.30%
Epoch: 2, Train Loss: 1.3206, Train Acc: 64.77%
Epoch: 2, Val Loss: 1.3412, Val Acc: 65.75%
Epoch: 3, Train Loss: 1.1885, Train Acc: 71.39%
Epoch: 3, Val Loss: 1.3391, Val Acc: 67.52%
Epoch: 4, Train Loss: 1.1152, Train Acc: 74.88%
Epoch: 4, Val Loss: 1.1641, Val Acc: 73.16%
Epoch: 5, Train Loss: 1.0671, Train Acc: 77.01%
Epoch: 5, Val Loss: 1.0711, Val Acc: 76.99%
Epoch: 6, Train Loss: 1.0100, Train Acc: 79.83%
Epoch: 6, Val Loss: 1.0254, Val Acc: 79.02%
Epoch: 7, Train Loss: 0.9786, Train Acc: 80.88%
Epoch: 7, Val Loss: 1.0319, Val Acc: 79.52%
Epoch: 8, Train Loss: 0.9464, Train Acc: 82.58%
Epoch: 8, Val Loss: 1.0027, Val Acc: 80.15%
Epoch: 9, Train Loss: 0.9225, Train Acc: 83.54%
Epoch: 9, Val Loss: 0.9849, Val Acc: 81.05%
Epoch: 10, Train Loss: 0.9066, Train Acc: 84.23%
Epoch: 10, Val Loss: 0.9103, Val Acc: 84.37%
Epoch: 11, Train Loss: 0.8877, Train Acc: 85.11%
Epoch: 11, Val Loss: 0.9673, Val Acc: 82.28%
Epoch: 12, Train Loss: 0.8759, Train Acc: 85.84%
Epoch: 12, Val Loss: 0.8815, Val Acc: 85.62%
Epoch: 13, Train Loss: 0.8636, Train Acc: 86.14%
Epoch: 13, Val Loss: 0.9741, Val Acc: 81.25%
Epoch: 14, Train Loss: 0.8557, Train Acc: 86.61%
Epoch: 14, Val Loss: 0.8411, Val Acc: 87.51%
Epoch: 15, Train Loss: 0.8437, Train Acc: 87.12%
Epoch: 15, Val Loss: 0.8641, Val Acc: 86.05%
Epoch: 16, Train Loss: 0.8336, Train Acc: 87.54%
Epoch: 16, Val Loss: 0.9138, Val Acc: 85.02%

```

Epoch: 17, Train Loss: 0.8222, Train Acc: 88.12%
Epoch: 17, Val Loss: 0.9515, Val Acc: 82.91%
Epoch: 18, Train Loss: 0.8146, Train Acc: 88.28%
Epoch: 18, Val Loss: 0.8485, Val Acc: 87.18%
Epoch: 19, Train Loss: 0.8007, Train Acc: 88.83%
Epoch: 19, Val Loss: 0.8211, Val Acc: 88.36%
Epoch: 20, Train Loss: 0.7938, Train Acc: 89.21%
Epoch: 20, Val Loss: 0.8512, Val Acc: 87.37%
Epoch: 21, Train Loss: 0.7869, Train Acc: 89.49%
Epoch: 21, Val Loss: 0.8446, Val Acc: 87.46%
Epoch: 22, Train Loss: 0.7735, Train Acc: 90.11%
Epoch: 22, Val Loss: 0.8370, Val Acc: 87.36%
Epoch: 23, Train Loss: 0.7656, Train Acc: 90.47%
Epoch: 23, Val Loss: 0.8504, Val Acc: 87.40%
Epoch: 24, Train Loss: 0.7571, Train Acc: 90.79%
Epoch: 24, Val Loss: 0.8313, Val Acc: 88.54%
Epoch: 25, Train Loss: 0.7467, Train Acc: 91.23%
Epoch: 25, Val Loss: 0.7748, Val Acc: 90.71%
Epoch: 26, Train Loss: 0.7399, Train Acc: 91.58%
Epoch: 26, Val Loss: 0.7636, Val Acc: 90.94%
Epoch: 27, Train Loss: 0.7319, Train Acc: 91.96%
Epoch: 27, Val Loss: 0.7966, Val Acc: 89.36%
Epoch: 28, Train Loss: 0.7247, Train Acc: 92.19%
Epoch: 28, Val Loss: 0.7617, Val Acc: 90.60%
Epoch: 29, Train Loss: 0.7116, Train Acc: 92.84%
Epoch: 29, Val Loss: 0.7628, Val Acc: 90.93%
Epoch: 30, Train Loss: 0.7053, Train Acc: 93.10%
Epoch: 30, Val Loss: 0.7528, Val Acc: 91.29%
Epoch: 31, Train Loss: 0.6968, Train Acc: 93.59%
Epoch: 31, Val Loss: 0.7558, Val Acc: 91.20%
Epoch: 32, Train Loss: 0.6907, Train Acc: 93.81%
Epoch: 32, Val Loss: 0.7402, Val Acc: 91.65%
Epoch: 33, Train Loss: 0.6835, Train Acc: 94.04%
Epoch: 33, Val Loss: 0.7385, Val Acc: 91.73%
Epoch: 34, Train Loss: 0.6737, Train Acc: 94.58%
Epoch: 34, Val Loss: 0.7403, Val Acc: 92.12%
Epoch: 35, Train Loss: 0.6670, Train Acc: 94.78%
Epoch: 35, Val Loss: 0.7476, Val Acc: 91.64%
Epoch: 36, Train Loss: 0.6584, Train Acc: 95.15%
Epoch: 36, Val Loss: 0.7368, Val Acc: 92.06%
Epoch: 37, Train Loss: 0.6522, Train Acc: 95.49%
Epoch: 37, Val Loss: 0.7236, Val Acc: 92.65%
Epoch: 38, Train Loss: 0.6444, Train Acc: 95.83%
Epoch: 38, Val Loss: 0.7290, Val Acc: 92.37%
Epoch: 39, Train Loss: 0.6366, Train Acc: 96.17%
Epoch: 39, Val Loss: 0.7251, Val Acc: 92.88%
Epoch: 40, Train Loss: 0.6277, Train Acc: 96.56%
Epoch: 40, Val Loss: 0.7298, Val Acc: 92.38%

Epoch: 41, Train Loss: 0.6234, Train Acc: 96.80%
Epoch: 41, Val Loss: 0.7277, Val Acc: 92.75%
Epoch: 42, Train Loss: 0.6155, Train Acc: 97.13%
Epoch: 42, Val Loss: 0.7126, Val Acc: 93.25%
Epoch: 43, Train Loss: 0.6092, Train Acc: 97.40%
Epoch: 43, Val Loss: 0.7132, Val Acc: 93.16%
Epoch: 44, Train Loss: 0.6046, Train Acc: 97.60%
Epoch: 44, Val Loss: 0.7106, Val Acc: 93.16%
Epoch: 45, Train Loss: 0.5986, Train Acc: 97.90%
Epoch: 45, Val Loss: 0.7081, Val Acc: 93.53%
Epoch: 46, Train Loss: 0.5937, Train Acc: 98.05%
Epoch: 46, Val Loss: 0.6984, Val Acc: 93.87%
Epoch: 47, Train Loss: 0.5899, Train Acc: 98.28%
Epoch: 47, Val Loss: 0.7031, Val Acc: 93.68%
Epoch: 48, Train Loss: 0.5844, Train Acc: 98.52%
Epoch: 48, Val Loss: 0.6981, Val Acc: 93.97%
Epoch: 49, Train Loss: 0.5804, Train Acc: 98.63%
Epoch: 49, Val Loss: 0.6984, Val Acc: 93.92%
Epoch: 50, Train Loss: 0.5782, Train Acc: 98.80%
Epoch: 50, Val Loss: 0.6969, Val Acc: 93.84%
Epoch: 51, Train Loss: 0.5761, Train Acc: 98.84%
Epoch: 51, Val Loss: 0.6989, Val Acc: 93.89%
Epoch: 52, Train Loss: 0.5717, Train Acc: 99.01%
Epoch: 52, Val Loss: 0.6919, Val Acc: 94.07%
Epoch: 53, Train Loss: 0.5720, Train Acc: 99.03%
Epoch: 53, Val Loss: 0.6945, Val Acc: 93.99%
Epoch: 54, Train Loss: 0.5693, Train Acc: 99.17%
Epoch: 54, Val Loss: 0.6910, Val Acc: 94.13%
Epoch: 55, Train Loss: 0.5662, Train Acc: 99.32%
Epoch: 55, Val Loss: 0.6892, Val Acc: 94.24%
Epoch: 56, Train Loss: 0.5669, Train Acc: 99.24%
Epoch: 56, Val Loss: 0.6882, Val Acc: 94.32%
Epoch: 57, Train Loss: 0.5649, Train Acc: 99.32%
Epoch: 57, Val Loss: 0.6884, Val Acc: 94.25%
Epoch: 58, Train Loss: 0.5651, Train Acc: 99.35%
Epoch: 58, Val Loss: 0.6881, Val Acc: 94.26%
Epoch: 59, Train Loss: 0.5652, Train Acc: 99.36%
Epoch: 59, Val Loss: 0.6888, Val Acc: 94.25%
Model saved to ./model_weights.pth

```
[10]: test_accuracy = evaluate(model, test_loader)
      print(f"Final Test Accuracy: {test_accuracy*100:.2f}%")
```

Final Test Accuracy: 94.25%

```
[11]: def imshow(img):
```



```

mean = np.array([0.4914, 0.4822, 0.4465])
std = np.array([0.2023, 0.1994, 0.2010])

img = img.numpy().transpose((1, 2, 0))
img = std * img + mean
img = np.clip(img, 0, 1)
plt.imshow(img)
plt.axis('off')

dataiter = iter(train_loader)
images, labels = next(dataiter)

classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', '
↳ 'ship', 'truck']

num_images = 16
rows = 4
cols = 4

fig, axs = plt.subplots(rows, cols, figsize=(10, 10))
axs = axs.ravel()

for i in range(num_images):
    axs[i].imshow(np.transpose((images[i] / 2 + 0.5).numpy(), (1, 2, 0)))
    axs[i].set_title(classes[labels[i]])
    axs[i].axis('off')
plt.tight_layout()
plt.show()

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).


WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).




```
[12]: from sklearn.metrics import confusion_matrix
import seaborn as sns


y_pred = []
y_true = []
model.eval()
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)
        y_pred.extend(predicted.view(-1).tolist())
        y_true.extend(labels.view(-1).tolist())

conf_matrix = confusion_matrix(y_true, y_pred)

plt.figure(figsize=(10, 8))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
            xticklabels=classes, yticklabels=classes)
plt.xlabel('Predicted Labels')
plt.ylabel('True Labels')
plt.title('Confusion Matrix')
plt.show()
```




```
[13]: plt.figure(figsize=(10, 4))
plt.plot(range(epochs), lr_changes, marker='o')
plt.title('Learning Rate Schedule')
plt.xlabel('Epoch')
plt.ylabel('Learning Rate')
plt.grid(True)
plt.show()
```



```
[14]: plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(train_accuracies, label='Training Accuracy')
plt.plot(val_accuracies, label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```



```
[15]: from torchsummary import summary
      print(summary(model, (3, 32, 32)))
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 66, 32, 32]	1,782
BatchNorm2d-2	[-1, 66, 32, 32]	132
Conv2d-3	[-1, 66, 32, 32]	39,204
BatchNorm2d-4	[-1, 66, 32, 32]	132
Conv2d-5	[-1, 66, 32, 32]	39,204
BatchNorm2d-6	[-1, 66, 32, 32]	132
BasicBlock-7	[-1, 66, 32, 32]	0
Conv2d-8	[-1, 66, 32, 32]	39,204
BatchNorm2d-9	[-1, 66, 32, 32]	132
Conv2d-10	[-1, 66, 32, 32]	39,204
BatchNorm2d-11	[-1, 66, 32, 32]	132
BasicBlock-12	[-1, 66, 32, 32]	0
Conv2d-13	[-1, 66, 32, 32]	39,204
BatchNorm2d-14	[-1, 66, 32, 32]	132
Conv2d-15	[-1, 66, 32, 32]	39,204
BatchNorm2d-16	[-1, 66, 32, 32]	132
BasicBlock-17	[-1, 66, 32, 32]	0
Conv2d-18	[-1, 132, 16, 16]	78,408
BatchNorm2d-19	[-1, 132, 16, 16]	264
Conv2d-20	[-1, 132, 16, 16]	156,816

BatchNorm2d-21	[-1, 132, 16, 16]	264
Conv2d-22	[-1, 132, 16, 16]	8,712
BatchNorm2d-23	[-1, 132, 16, 16]	264
BasicBlock-24	[-1, 132, 16, 16]	0
Conv2d-25	[-1, 132, 16, 16]	156,816
BatchNorm2d-26	[-1, 132, 16, 16]	264
Conv2d-27	[-1, 132, 16, 16]	156,816
BatchNorm2d-28	[-1, 132, 16, 16]	264
BasicBlock-29	[-1, 132, 16, 16]	0
Conv2d-30	[-1, 132, 16, 16]	156,816
BatchNorm2d-31	[-1, 132, 16, 16]	264
Conv2d-32	[-1, 132, 16, 16]	156,816
BatchNorm2d-33	[-1, 132, 16, 16]	264
BasicBlock-34	[-1, 132, 16, 16]	0
Conv2d-35	[-1, 132, 16, 16]	156,816
BatchNorm2d-36	[-1, 132, 16, 16]	264
Conv2d-37	[-1, 132, 16, 16]	156,816
BatchNorm2d-38	[-1, 132, 16, 16]	264
BasicBlock-39	[-1, 132, 16, 16]	0
Conv2d-40	[-1, 264, 8, 8]	313,632
BatchNorm2d-41	[-1, 264, 8, 8]	528
Conv2d-42	[-1, 264, 8, 8]	627,264
BatchNorm2d-43	[-1, 264, 8, 8]	528
Conv2d-44	[-1, 264, 8, 8]	34,848
BatchNorm2d-45	[-1, 264, 8, 8]	528
BasicBlock-46	[-1, 264, 8, 8]	0
Conv2d-47	[-1, 264, 8, 8]	627,264
BatchNorm2d-48	[-1, 264, 8, 8]	528
Conv2d-49	[-1, 264, 8, 8]	627,264
BatchNorm2d-50	[-1, 264, 8, 8]	528
BasicBlock-51	[-1, 264, 8, 8]	0
Conv2d-52	[-1, 264, 8, 8]	627,264
BatchNorm2d-53	[-1, 264, 8, 8]	528
Conv2d-54	[-1, 264, 8, 8]	627,264
BatchNorm2d-55	[-1, 264, 8, 8]	528
BasicBlock-56	[-1, 264, 8, 8]	0
AdaptiveAvgPool2d-57	[-1, 264, 1, 1]	0
Linear-58	[-1, 10]	2,650

=====

Total params: 4,916,284

Trainable params: 4,916,284

Non-trainable params: 0

Input size (MB): 0.01

Forward/backward pass size (MB): 16.63

Params size (MB): 18.75

Estimated Total Size (MB): 35.40

None

```
[3]: import sys
import torch

!cat /proc/cpuinfo | grep 'model name' | uniq | awk -F: '{print "CPU:", $2}'
!nvidia-smi --query-gpu=gpu_name --format=csv,noheader | awk '{$1=$1};1' | awk
↪ '{print "GPU:", $0}'
!cat /proc/meminfo | grep 'MemTotal' | awk '{print "System Memory:", $2/1024/
↪ 1024, "GB"}'
print("Python Version:", sys.version)
print("CUDA version:", torch.version.cuda)
print("Torch Version:", torch.__version__)
```

CPU: Intel(R) Xeon(R) CPU @ 2.20GHz

GPU: Tesla V100-SXM2-16GB

System Memory: 50.9937 GB

Python Version: 3.10.12 (main, Nov 20 2023, 15:14:05) [GCC 11.4.0]

CUDA version: 12.1

Torch Version: 2.2.1+cu121